

```
In [3]: import pandas as pd
import os
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
```

```
In [4]: #Changing Directory
os.chdir("C:\\Users\\rakbansal\\Desktop\\Kaggle\\Earthquake")
```

```
In [5]: #Reading in the files
train_data= pd.read_csv("train_values.csv", index_col= 'building_id')
train_labels= pd.read_csv("train_labels.csv", index_col= 'building_id')
test_data= pd.read_csv("test_values.csv", index_col= "building_id")
```

```
In [6]: # Checking shape and size of the data
print('training data shape: ', train_data.shape, '\ntest data shape: ', te
st_data.shape, '\nTraining Label Shape: ', train_labels.shape)
```

```
training data shape: (260601, 38)
test data shape: (86868, 38)
Training Label Shape: (260601, 1)
```

```
In [7]: #Let's take a look at first few values of training data
train_data.head()
```

Out[7]:

|             | geo_level_1_id | geo_level_2_id | geo_level_3_id | count_floors_pre_eq | age | area_percentage |
|-------------|----------------|----------------|----------------|---------------------|-----|-----------------|
| building_id |                |                |                |                     |     |                 |
| 802906      | 6              | 487            | 12198          | 2                   | 30  | 6               |
| 28830       | 8              | 900            | 2812           | 2                   | 10  | 8               |
| 94947       | 21             | 363            | 8973           | 2                   | 10  | 5               |
| 590882      | 22             | 418            | 10694          | 2                   | 10  | 6               |
| 201944      | 11             | 131            | 1488           | 3                   | 30  | 8               |

5 rows × 38 columns

```
In [8]: #Let's take a look at our labels
train_labels.head()
```

Out[8]:

|             | damage_grade |
|-------------|--------------|
| building_id |              |
| 802906      | 3            |
| 28830       | 2            |
|             |              |

|               |          |
|---------------|----------|
| <b>94947</b>  | <b>3</b> |
| <b>590882</b> | <b>2</b> |
| <b>201944</b> | <b>3</b> |

```
In [4]: #Defining train and label
train_label= train_labels[['damage_grade']]
```

```
In [9]: #Let's take a look if Training Data has any missing values
train_data.isnull().sum()
```

```
Out[9]: geo_level_1_id          0
geo_level_2_id          0
geo_level_3_id          0
count_floors_pre_eq     0
age                     0
area_percentage         0
height_percentage       0
land_surface_condition  0
foundation_type         0
roof_type               0
ground_floor_type       0
other_floor_type        0
position                0
plan_configuration      0
has_superstructure_adobe_mud          0
has_superstructure_mud_mortar_stone  0
has_superstructure_stone_flag        0
has_superstructure_cement_mortar_stone 0
has_superstructure_mud_mortar_brick  0
has_superstructure_cement_mortar_brick 0
has_superstructure_timber             0
has_superstructure_bamboo             0
has_superstructure_rc_non_engineered  0
has_superstructure_rc_engineered     0
has_superstructure_other              0
legal_ownership_status               0
count_families                      0
has_secondary_use                    0
has_secondary_use_agriculture         0
has_secondary_use_hotel               0
has_secondary_use_rental              0
has_secondary_use_institution         0
has_secondary_use_school              0
has_secondary_use_industry            0
has_secondary_use_health_post         0
has_secondary_use_gov_office          0
has_secondary_use_use_police          0
has_secondary_use_other               0
dtype: int64
```

```
In [10]: #Let's check the same for test values
test_data.isnull().sum()
```

```
Out[10]: geo_level_1_id          0
geo_level_2_id          0
```

```

geo_level_3_id          0
count_floors_pre_eq     0
age                     0
area_percentage         0
height_percentage       0
land_surface_condition  0
foundation_type         0
roof_type               0
ground_floor_type       0
other_floor_type        0
position                0
plan_configuration      0
has_superstructure_adobe_mud          0
has_superstructure_mud_mortar_stone  0
has_superstructure_stone_flag        0
has_superstructure_cement_mortar_stone 0
has_superstructure_mud_mortar_brick  0
has_superstructure_cement_mortar_brick 0
has_superstructure_timber             0
has_superstructure_bamboo             0
has_superstructure_rc_non_engineered  0
has_superstructure_rc_engineered      0
has_superstructure_other              0
legal_ownership_status               0
count_families                      0
has_secondary_use                    0
has_secondary_use_agriculture        0
has_secondary_use_hotel              0
has_secondary_use_rental             0
has_secondary_use_institution        0
has_secondary_use_school             0
has_secondary_use_industry           0
has_secondary_use_health_post        0
has_secondary_use_gov_office         0
has_secondary_use_use_police         0
has_secondary_use_other              0
dtype: int64

```

```
In [11]: train_labels.isnull().sum()
```

```
Out[11]: damage_grade      0
dtype: int64
```

Great! None of the dataset have any missing value. Let's now move on to Exploratory Data Analysis. First, let's see what kind of features do we have.

```
In [12]: train_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 260601 entries, 802906 to 747594
Data columns (total 38 columns):
geo_level_1_id          260601 non-null int64
geo_level_2_id          260601 non-null int64
geo_level_3_id          260601 non-null int64
count_floors_pre_eq     260601 non-null int64
age                     260601 non-null int64

```

```

area_percentage                260601 non-null int64
height_percentage              260601 non-null int64
land_surface_condition        260601 non-null object
foundation_type               260601 non-null object
roof_type                    260601 non-null object
ground_floor_type            260601 non-null object
other_floor_type             260601 non-null object
position                     260601 non-null object
plan_configuration            260601 non-null object
has_superstructure_adobe_mud  260601 non-null int64
has_superstructure_mud_mortar_stone 260601 non-null int64
has_superstructure_stone_flag 260601 non-null int64
has_superstructure_cement_mortar_stone 260601 non-null int64
has_superstructure_mud_mortar_brick 260601 non-null int64
has_superstructure_cement_mortar_brick 260601 non-null int64
has_superstructure_timber     260601 non-null int64
has_superstructure_bamboo     260601 non-null int64
has_superstructure_rc_non_engineered 260601 non-null int64
has_superstructure_rc_engineered 260601 non-null int64
has_superstructure_other      260601 non-null int64
legal_ownership_status        260601 non-null object
count_families                260601 non-null int64
has_secondary_use             260601 non-null int64
has_secondary_use_agriculture 260601 non-null int64
has_secondary_use_hotel       260601 non-null int64
has_secondary_use_rental      260601 non-null int64
has_secondary_use_institution 260601 non-null int64
has_secondary_use_school      260601 non-null int64
has_secondary_use_industry     260601 non-null int64
has_secondary_use_health_post 260601 non-null int64
has_secondary_use_gov_office  260601 non-null int64
has_secondary_use_use_police   260601 non-null int64
has_secondary_use_other       260601 non-null int64
dtypes: int64(30), object(8)
memory usage: 69.6+ MB

```

Most of our variables are factor or categorical variables. We will deal with them later. Let's separate the numerical variables first.

```

In [13]: numerical_features= ['count_families', 'count_floors_pre_eq', 'age', 'area_
percentage', 'height_percentage']
num_data= train_data[numerical_features]

```

```

In [15]: #Let's visualize these numerical features
import seaborn as sns

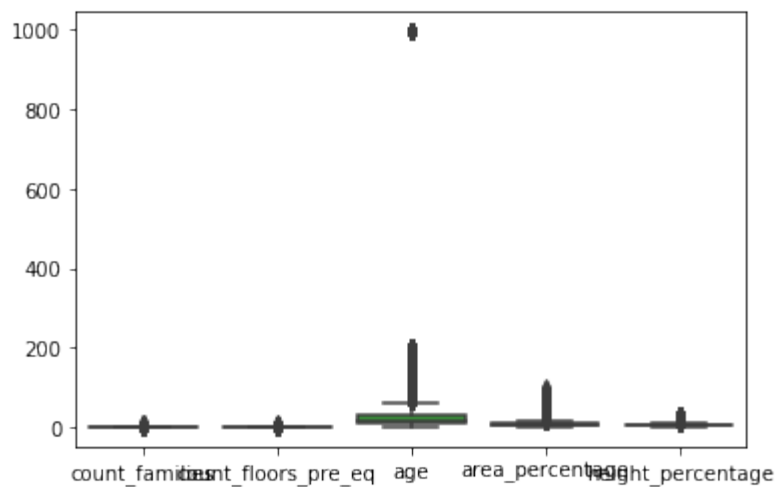
sns.boxplot(data=num_data)

```

```

Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0xd2e2f10>


```



We can see that most values are distributed from 0-200, except 'age' of the buildings which seems to have a few outliers. Let's take a look at it.


```
In [16]: sns.countplot(num_data['age'])
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x113afa70>
```



We can see that most of the buildings are new. However, there are couple that are over 995 years old. Let's take a closer look at all numerical features.

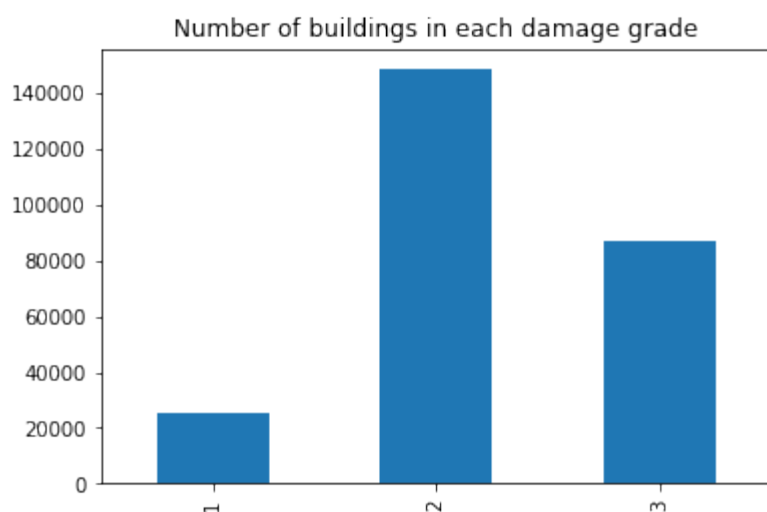
```
In [17]: #Let's now see the correlation of all our numerical features.
plt.figure(figsize=(15,10))
correlation = num_data.corr()
sns.heatmap(correlation, annot=True, cmap="YlGnBu")
plt.show();
```



We can see that there is no significant correlation among all our numerical variables except 'count\_floor\_per\_eq and 'height\_percentage'. We will take care of that during our feature selection.

```
In [18]: #lets's see our training data
(train_labels.damage_grade.value_counts().sort_index().plot.bar(title= 'Number of buildings in each damage grade'))
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0xd503670>
```



Looks like most buildings had medium damage.


```
In [19]: #Let's analyze our categorical variables
merged_data= pd.merge(train_data, train_labels, on= 'building_id')
# Set up a factorplot
g = sns.factorplot("foundation_type", "height_percentage", "damage_grade",
    data=merged_data, kind="bar", size=6, palette="muted", legend=False)
plt.show()
```

c:\users\rakbansal\appdata\local\programs\python\python37-32\lib\site-packages\seaborn\categorical.py:3666: UserWarning: The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed to `strip` in `catplot`.

warnings.warn(msg)

c:\users\rakbansal\appdata\local\programs\python\python37-32\lib\site-packages\seaborn\categorical.py:3672: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

warnings.warn(msg, UserWarning)



```
In [20]: #Let's do some feature selection. Let's run a Random Forest to get feature
importances.
train_label= train_labels[['damage_grade']]
```

```
In [21]: #Coding categorical Data
train_dummy= pd.get_dummies(train_data)
test_dummy= pd.get_dummies(test_data)
```

```
In [22]: #Creating a Random Forest Model
```

```
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor(n_estimators=10, random_state=5, max_depth=10)
```

```
In [23]: #Fitting the model on Test Data
model.fit(train_dummy,train_label)
```

```
c:\users\rakbansal\appdata\local\programs\python\python37-32\lib\site-packages\ipykernel_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
Out[23]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=10,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10,
                                n_jobs=None, oob_score=False, random_state=5, verbose=0,
                                warm_start=False)
```

Let's Take a look at 'Feature Importances' of our model based on which we will select the most significant features.

```
In [29]: #Feature importance in Descending order numbers
feature_importances = pd.Series(model.feature_importances_, index= train_dummy.columns).sort_values(ascending=False)
print(importances)
```

|  |          |
|--|----------|
| geo_level_1_id                         | 0.435828 |
| foundation_type_r                      | 0.302499 |
| has_superstructure_mud_mortar_stone    | 0.063477 |
| geo_level_2_id                         | 0.056666 |
| age                                    | 0.021128 |
| foundation_type_i                      | 0.017616 |
| has_superstructure_stone_flag          | 0.013464 |
| has_superstructure_cement_mortar_brick | 0.013269 |
| geo_level_3_id                         | 0.012858 |
| other_floor_type_q                     | 0.005916 |
| area_percentage                        | 0.005730 |
| count_families                         | 0.005020 |
| height_percentage                      | 0.004641 |
| roof_type_q                            | 0.004253 |
| has_superstructure_mud_mortar_brick    | 0.003993 |
| roof_type_n                            | 0.003654 |
| has_superstructure_timber              | 0.002421 |
| count_floors_pre_eq                    | 0.002092 |
| has_secondary_use                      | 0.001867 |
| foundation_type_h                      | 0.001857 |
| roof_type_x                            | 0.001836 |
| has_superstructure_adobe_mud           | 0.001509 |
| other_floor_type_x                     | 0.001503 |
| ground_floor_type_v                    | 0.001450 |
| has_superstructure_rc_non_engineered   | 0.001407 |
| ground_floor_type_f                    | 0.001140 |



```

has_superstructure_other          0.001038
foundation_type_u                 0.001028
position_s                        0.001024
has_superstructure_cement_mortar_stone 0.000843
...
other_floor_type_j                0.000386
other_floor_type_s                0.000364
has_superstructure_rc_engineered  0.000352
legal_ownership_status_v          0.000340
position_j                        0.000340
position_t                        0.000270
plan_configuration_d              0.000244
legal_ownership_status_a          0.000236
ground_floor_type_m              0.000218
ground_floor_type_z              0.000206
legal_ownership_status_r          0.000196
foundation_type_w                 0.000185
plan_configuration_q              0.000184
plan_configuration_u              0.000167
legal_ownership_status_w          0.000139
has_secondary_use_industry        0.000138
has_secondary_use_rental          0.000118
position_o                        0.000041
plan_configuration_a              0.000027
plan_configuration_s              0.000026
has_secondary_use_institution     0.000015
has_secondary_use_school          0.000010
has_secondary_use_health_post     0.000008
has_secondary_use_gov_office      0.000003
plan_configuration_o              0.000002
plan_configuration_c              0.000001
plan_configuration_f              0.000000
plan_configuration_m              0.000000
plan_configuration_n              0.000000
has_secondary_use_use_police       0.000000
Length: 68, dtype: float64

```

```

In [35]: #Let's Visualize the same
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# Creating a bar plot
feature_importances= feature_importances.nlargest(25) #Selects tehe top 25
features
sns.barplot(x=feature_importances, y=feature_importances.index)
# Add labels to your graph

plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()


```

No handles with labels found to put in legend.

```

Out[35]: <matplotlib.legend.Legend at 0x1b0025b0>

```



```
In [39]: #Creating a list of our newly selected top 25 Features
new_features= list(feature_importances.index)
```

```
In [40]: #Initializing our training and testing data with new features
training_dataset= train_dummy[new_features]
testing_dataset= test_dummy[new_features]
```

```
In [43]: #Let's take a look at new training dataset
training_dataset.head()
```

Out[43]:

|             | geo_level_1_id | foundation_type_r | has_superstructure_mud_mortar_stone | geo_level_2_id |
|-------------|----------------|-------------------|-------------------------------------|----------------|
| building_id |                |                   |                                     |                |
| 802906      | 6              | 1                 | 1                                   | 487            |
| 28830       | 8              | 1                 | 1                                   | 900            |
| 94947       | 21             | 1                 | 1                                   | 363            |
| 590882      | 22             | 1                 | 1                                   | 418            |
| 201944      | 11             | 1                 | 0                                   | 131            |

5 rows × 5 columns

Let us now divide the training dataset into 'Training data' and 'validation data' before we apply it on the 'test data'.

```
In [49]: #Dividing training data into train and validation dataset
from sklearn.model_selection import train_test_split
X= training_dataset
y= train_label
```

```
In [50]: # Split dataset into training set and validation set
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3) # 7
```

0% training and 30% test

```
In [70]: # Create the model with 200 trees
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=200, oob_score = True, n_jobs
= -1, random_state =10, max_features = "auto", min_samples_leaf = 50)
```

```
In [71]: #Let's fit the model on training dataset
trained=model.fit(X_train, y_train)

c:\users\rakbansal\appdata\local\programs\python\python37-32\lib\site-pack
ages\ipykernel_launcher.py:2: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to (n_s
amples,), for example using ravel().
```

```
In [72]: #Predict on Validation dataset
predicted= model.predict(X_val)
```

```
In [63]: #Test the accuracy of the Model- Import required libraries
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn import metrics
```

```
In [73]: print("Accuracy:",metrics.accuracy_score(y_val, predicted))

Accuracy: 0.6918049142374746
```

```
In [56]: from sklearn.metrics import classification_report
```

```
In [74]: print (classification_report(y_val, predicted) )
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1            | 0.66      | 0.34   | 0.45     | 7546    |
| 2            | 0.68      | 0.89   | 0.77     | 44580   |
| 3            | 0.75      | 0.46   | 0.57     | 26055   |
| accuracy     |           |        | 0.69     | 78181   |
| macro avg    | 0.70      | 0.56   | 0.60     | 78181   |
| weighted avg | 0.70      | 0.69   | 0.67     | 78181   |

The model provides an accuracy of 0.69

```
In [75]: #Let's make predictions on the test dataset
damage_predicted= model.predict(testing_dataset)
```

```
In [78]: print(damage_predicted)

[3 2 2 ... 2 2 2]
```

```
In [80]: damage_predicted.shape
```

Out[80]: (86868,)

In [ ]: